

MSDS 697

Distributed Data Systems

DIANE WOODBRIDGE, PH.D



Week 3 - Comments (What you liked/disliked so far? What should I do for you?)

Alumni Feedback

Monday, November 1st



Shruti Roy 9:00 AM

Hi professor Diane,
Happy Halloween 🎃 to you and your family!

I just wanted to say thank you for your distributed computing class! I am setting up a lot of pipelines while working at RDC and I was genuinely able to do that a lot because of what I learnt in your classes.



dwoodbridge 9:00 AM

Awww - this means so much to me! I will share what you said in my class 😊 I am teaching it this module! 😊



1



Shruti Roy 9:01 AM

Absolutely Professor Diane! Super useful stuff



1



9:01 I hope you and your family are in good health and doing well ❤️



UNIVERSITY OF SAN FRANCISCO

CHANGE THE WORLD FROM HERE

Announcement

Group Assignment 1 - Due tomorrow midnight!

- Task 1 (Nov 5) : Data Selection

Quiz - 1st floor

- Nov 11, **9:00 - 9:55 AM** : Multiple Choice (Closed-book), and Programming (Open-book).
- Study list and example quiz are available.

Office Hour : Tuesday 9:15-10 am

Professionalism

- **Class Attendance**

- No cellphones, social media, slack, texting during the class .
- Please only use laptops for class-related purposes.

- **Assignment**

- No late submissions are allowed.
- Do not share any homework and exam files - All the codes including the last 5 years will be tested by Moss (Measure Of Software Similarity).
- Make sure that your code runs in Python 3.10 and Pyspark 3.3

Review

MapReduce

Hadoop MapReduce vs. Spark

Spark Overview

Spark Components

Spark Examples

Resilient Distributed Dataset (RDD)

Running a Spark application (.py)
using the spark-submit

RDD Operations

- Transformation
- Action (Week 3)



Contents

RDD Operations

- Transformation
- Action



RDD Operations - Transformation

Transformation

- Construct a new RDD from an existing RDD.

```
line_with_spark = lines.filter(lambda x : "spark" in x)
```

- Lazy Evaluation.
 - Computation doesn't take place until an action is triggered.
- Return new RDDs.
 - Doesn't change the original RDDs.



RDD Operations - Transformation

Transformation Operation Types

<code>map(func)</code>	Return a new distributed dataset formed by passing each element of the source through a function <i>func</i> .
<code>filter(func)</code>	Return a new dataset formed by selecting those elements of the source on which <i>func</i> returns true.
<code>distinct()</code>	Return a new dataset that contains the distinct elements of the source dataset.
<code>union(otherDataset)</code>	Return a new dataset that contains the union of the elements in the source dataset and the argument.
<code>intersection(otherDataset)</code>	Return a new RDD that contains the intersection of elements in the source dataset and the argument.
<code>subtract(otherDataset)</code>	Return each value in self that is not contained in otherDataset.
<code>cartesian(otherDataset)</code>	Return the Cartesian product of the RDD and otherDataset.
<code>sortBy(func, ascending=True)</code>	Sorts this RDD by the given <i>func</i> .

<http://spark.apache.org/docs/latest/programming-guide.html>



Contents

RDD Operations

- Transformation
- **Action**



RDD Operations - Action

- Compute a result based on an RDD and return the result to the driver program.

```
input_rdd.count()  
input_rdd.collect()
```

- Return non-RDDs like a single number, string, array, etc.



Today's Example

For data/numbers.txt, calculate the (count, sum) of all odd numbers.

Expected output : (49, 233)

4
0
3
6
8
7
9
9
8
5
6
5
6
3
6
5
1
8
8
1
5
4
1
5
8
2
2
6
8
8
3
9
9
0
7
3
6



RDD Operations - Action

Action Operation Types

<code>collect()</code>	Return all the elements of the dataset as an array at the driver program. This is usually useful after a filter or other operation that returns a sufficiently small subset of the data.
<code>first()</code>	Return the first element of the dataset (similar to <code>take(1)</code>).
<code>take(n)</code>	Return an array with the first <code>n</code> elements of the dataset.
<code>count()</code>	Return the number of elements in the dataset.
<code>countByValue()</code>	Return the number of times each element occurs in the RDD.
<code>reduce(func)</code>	Combine the elements of the RDD together in parallel. (eg. Sum)
<code>fold(zero, func)</code>	Same as <code>reduce()</code> , but with the provided zero value.
<code>aggregate(zero, SeqOp, combOp)</code>	Similar to <code>reduce()</code> but used to return a different type.
<code>saveAsTextFile(path)</code>	Write the elements of the dataset as a text file (or set of text files) in a given directory in the local filesystem, HDFS or any other Hadoop-supported file system. Spark will call <code>toString</code> on each element to convert it to a line of text in the file.
<code>mean()</code> , <code>sum()</code> , <code>min()</code> , <code>max()</code> , <code>variance()</code> , <code>stdev()</code>	Returns mean, sum, minimum, maximum, variance and standard deviation.

RDD Operation - Action

Numeric RDD Action Types

- `mean()`
 - Return the mean of the RDD's elements.
- `sum()`
 - Add up the elements in the RDD.
- `max()`
 - Return the maximum item in the RDD.
- `min()`
 - Return the minimum item in the RDD.
- `variance()`
 - Return the variance of the RDD's elements.
- `stdev()`
 - Return the standard deviation of the RDD's elements.

<https://spark.apache.org/docs/latest/api/python/pyspark.html#pyspark.RDD>



ex01

For data/numbers.txt, calculate the count, sum, mean, max, min, variance and standard deviation for the odd numbers.

```
odd_numbers.count()
```

49

```
odd_numbers.sum()
```

233

```
odd_numbers.mean()
```

4.755102040816326

```
odd_numbers.max()
```

9

```
odd_numbers.min()
```

1

```
odd_numbers.variance()
```

8.10329029571012

```
odd_numbers.stdev()
```

2.8466278814959503



RDD Operations - Action

`collect()`

- Return all the elements of the dataset as an array at the driver program.
- This is usually useful after a filter or other operation that returns a sufficiently small subset of the data.

`first()`

- Return the first element of the dataset (similar to `take(1)`).

`take(n)`

- Return an array with the first n elements of the dataset.

`count()`

- Return the number of elements in the dataset.

`countByValue()`

- Return the number of times each element occurs in the RDD.

ex02

For data/numbers.txt,

- What are the odd numbers in the file?
- What is the first odd number?
- What are the 5 odd numbers?
- What are the occurrences of each odd number?

<https://spark.apache.org/docs/latest/api/python/pyspark.html#pyspark.RDD>



ex02

For data/numbers.txt,

- What are the odd numbers in the file?
- What is the first odd number?
- What are the 5 odd numbers?
- What are the occurrences of each odd number?

```
odd_numbers.collect()
```

```
odd_numbers.first()
```

```
odd_numbers.take(5)
```

```
odd_numbers.count()
```

```
odd_numbers.countByValue()
```

```
odd_numbers.countByValue().items()
```





`rdd.count().count()` works without any errors, where rdd is an RDD.

True

False



`rdd.count().count()` works without any errors, where `rdd` is an RDD.

True

False



`rdd.count().count()` works without any errors, where `rdd` is an RDD.

True

False

RDD Operations - Action

reduce(func)

- Take a function that operates on **two** elements of the type in your RDD.
- Returns a new element of the same type.
- Ex. `sum = rdd.reduce(lambda x, y : x+y)`

<https://spark.apache.org/docs/latest/api/python/pyspark.html#pyspark.RDD>



ex03

For data/numbers.txt, calculate the sum of the odd numbers.

ex03

For data/numbers.txt, calculate the sum of the odd numbers.

```
numbers = sc.textFile("../data/numbers.txt")
```

```
odd_numbers = numbers.map(lambda x : int(x))\  
                      .filter(lambda x : x % 2 == 1)
```

```
odd_numbers.reduce(lambda x,y : x + y)
```



RDD Operations - Action

`fold(zeroValue, func)`

- Take a function with the same signature as needed for `reduce()`.
- Take a “zero value” to be used for the initial call on each partition.
- Returns a new element of the same type.
- `fold()` is useful when the given RDD is empty. cf. `reduce()`

<https://spark.apache.org/docs/latest/api/python/pyspark.html#pyspark.RDD>



ex04

For data/numbers.txt, calculate the sum of the odd numbers with different zero values using fold().

ex04

For data/numbers.txt, calculate the sum of the odd numbers with different zero values using fold().

```
numbers.getNumPartitions()
```

8

```
odd_numbers = numbers.map(lambda x : int(x))\  
                    .filter(lambda x : x%2 == 1)
```

```
odd_numbers.fold(0, lambda x,y : x+y)
```

233

```
odd_numbers.fold(1, lambda x,y : x+y)
```

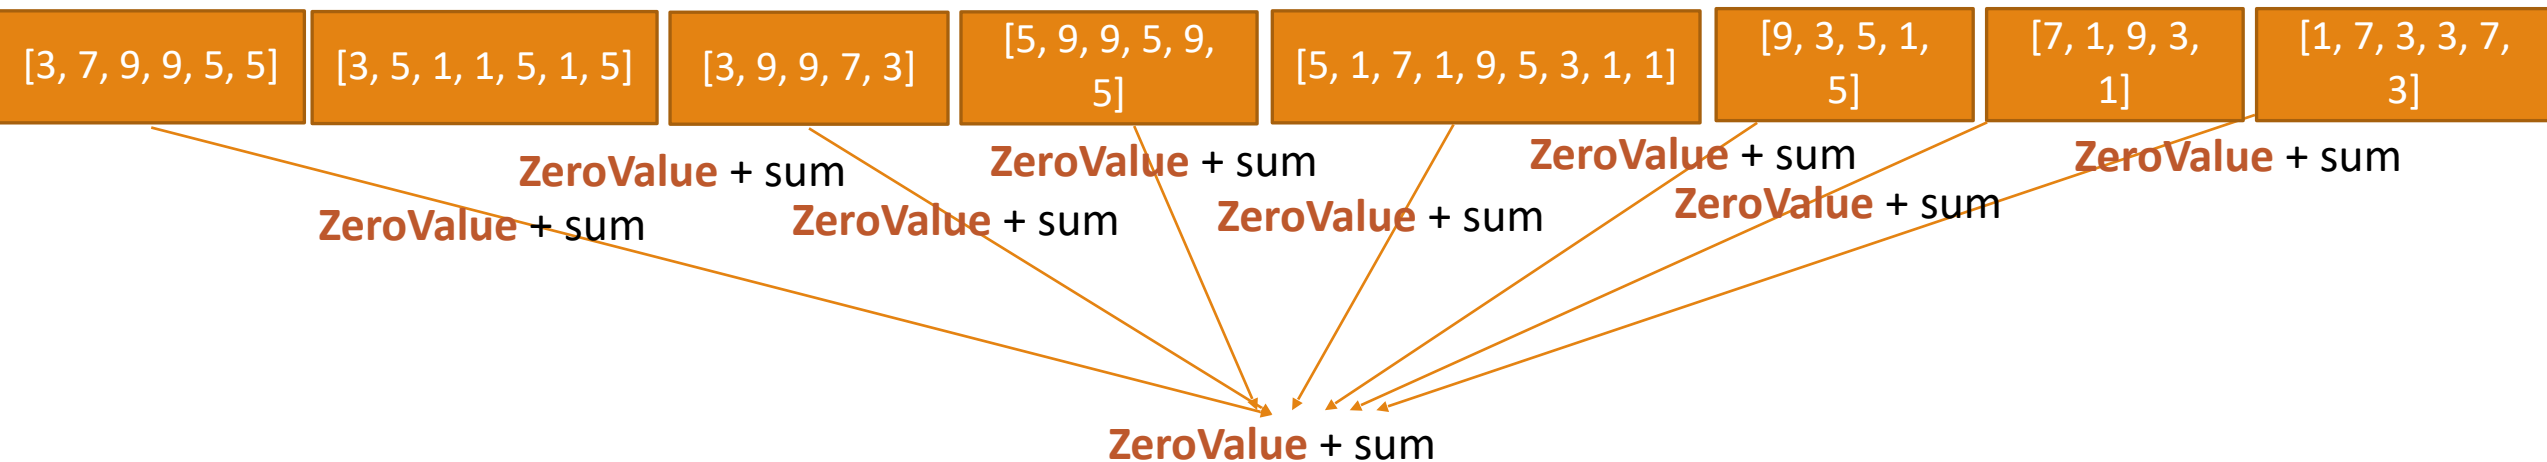
242

```
odd_numbers.fold(2, lambda x,y : x+y)
```

251

ex04

For data/numbers.txt, calculate the sum of the odd numbers with different zero values using fold().



ex04

For data/numbers.txt, calculate the sum of the odd numbers with different zero values using fold().

To do : Try it with different number partitions.



ex04

For data/numbers.txt, calculate the sum of the odd numbers with different zero values using fold().

fold() is useful when the given RDD is empty. cf. reduce()

```
numbers = sc.parallelize([])
```

```
numbers.reduce(lambda x,y : x+y)
```

```
-----  
ValueError                                Traceback (most recent call last)  
<ipython-input-19-0efd3e2d4ec0> in <module>  
----> 1 numbers.reduce(lambda x,y : x+y)  
  
~/anaconda3/envs/MSDS694/lib/python3.6/site-packages/pyspark/rdd.py in reduce(self, f)  
    863         if vals:  
    864             return reduce(f, vals)  
--> 865         raise ValueError("Can not reduce() empty RDD")  
    866  
    867     def treeReduce(self, f, depth=2):  
  
ValueError: Can not reduce() empty RDD
```

```
numbers.fold(0, lambda x,y : x+y)
```

RDD Operations - Action

`aggregate(zeroValue, seqOp, combOp)`

- Aggregate the elements of each partition using `zeroValue` and `seqOp` → Aggregate the results for all the partitions, using `combOp` and `zeroValue`.
- `zeroValue` : should be in a format we want to return.
 - Can return an element of a different type from the original RDD.
- `seqOp` : Function to combine the elements from the RDD. Runs in a partition.
- `combOp` : Function to merge two accumulators resulted from `seqOp` and `zeroValue`, given that each nodes accumulates its own results locally.

https://spark.apache.org/docs/latest/api/python/_modules/pyspark/rdd.html#RDD.aggregate



ex05

For data/numbers.txt, calculate the number (count) and sum of the odd numbers as a pair using the aggregate function.



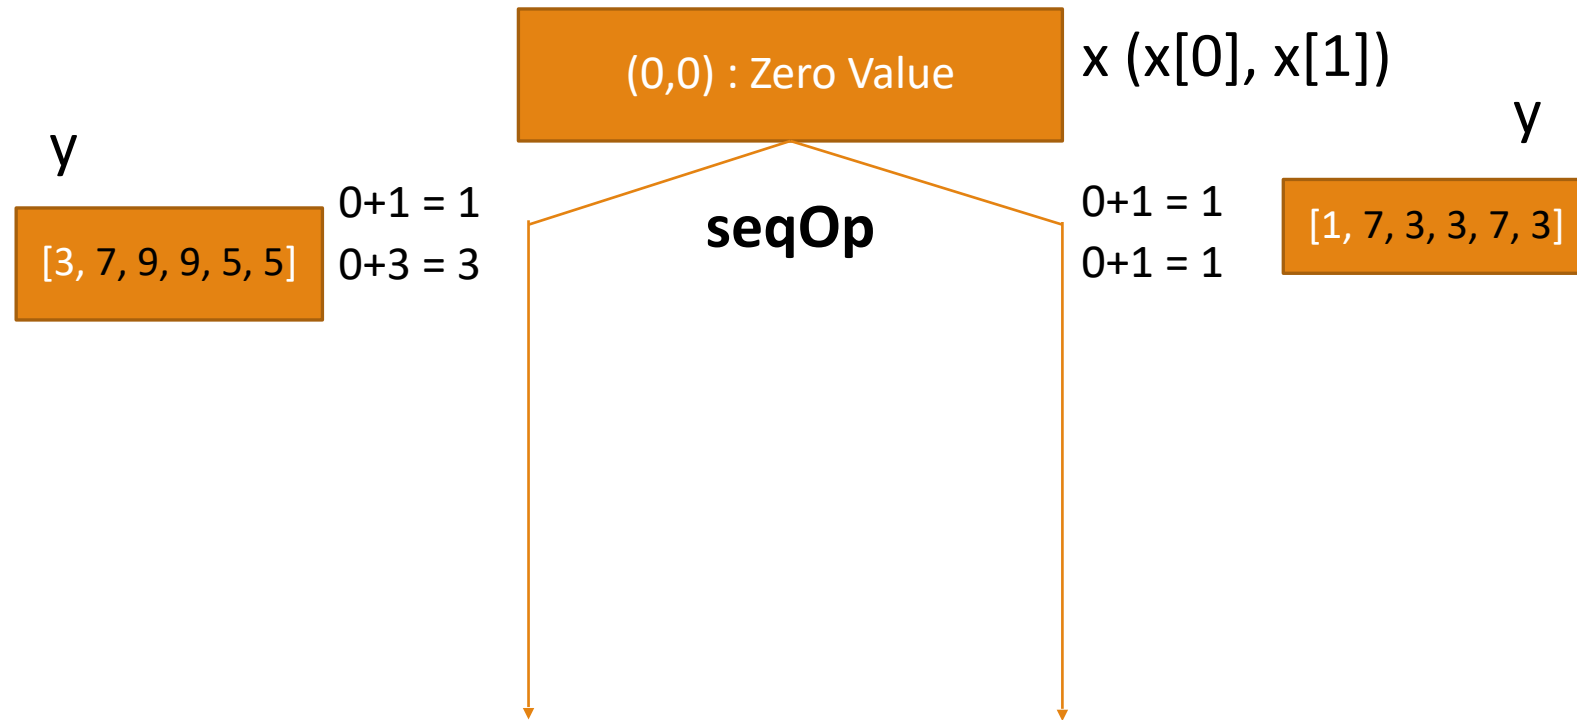
ex05

For data/numbers.txt, calculate the number (count) and sum of the odd numbers as a pair using the aggregate function.

- The type to return : (Count, Sum) → (Int, Int) - the format of the zero value
- In the same partition (like Map)
 - Every time encountering a new number, increase the counter by 1
 - counter + 1
 - Every time encountering a new number, add the two numbers
 - accumulated sum + new number
- Merging results from partitions (like the last step in Reduce)
 - Add resulted counters together
 - Add resulted sums together

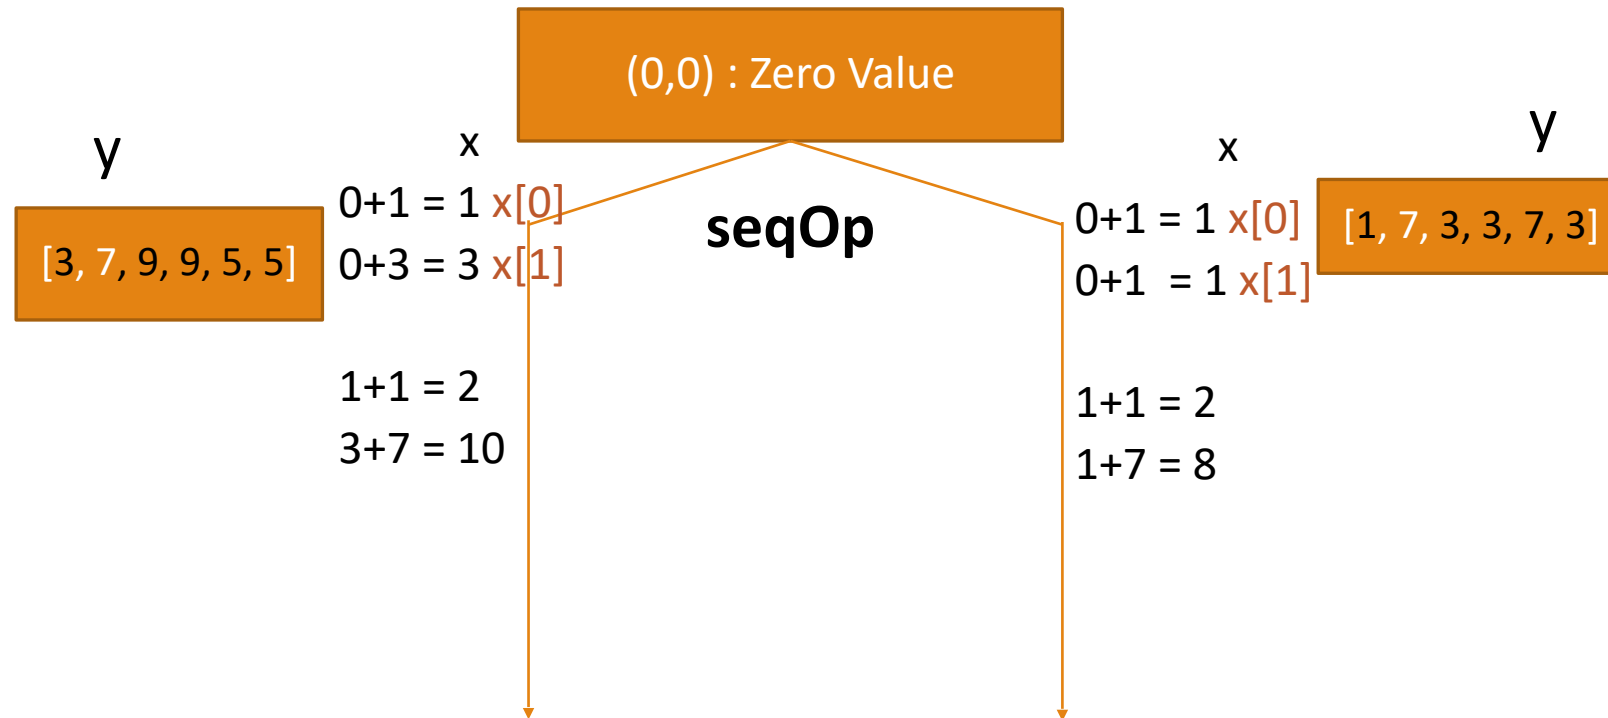
ex05

For data/numbers.txt, calculate the number and sum of the odd numbers as a pair using the aggregate function.



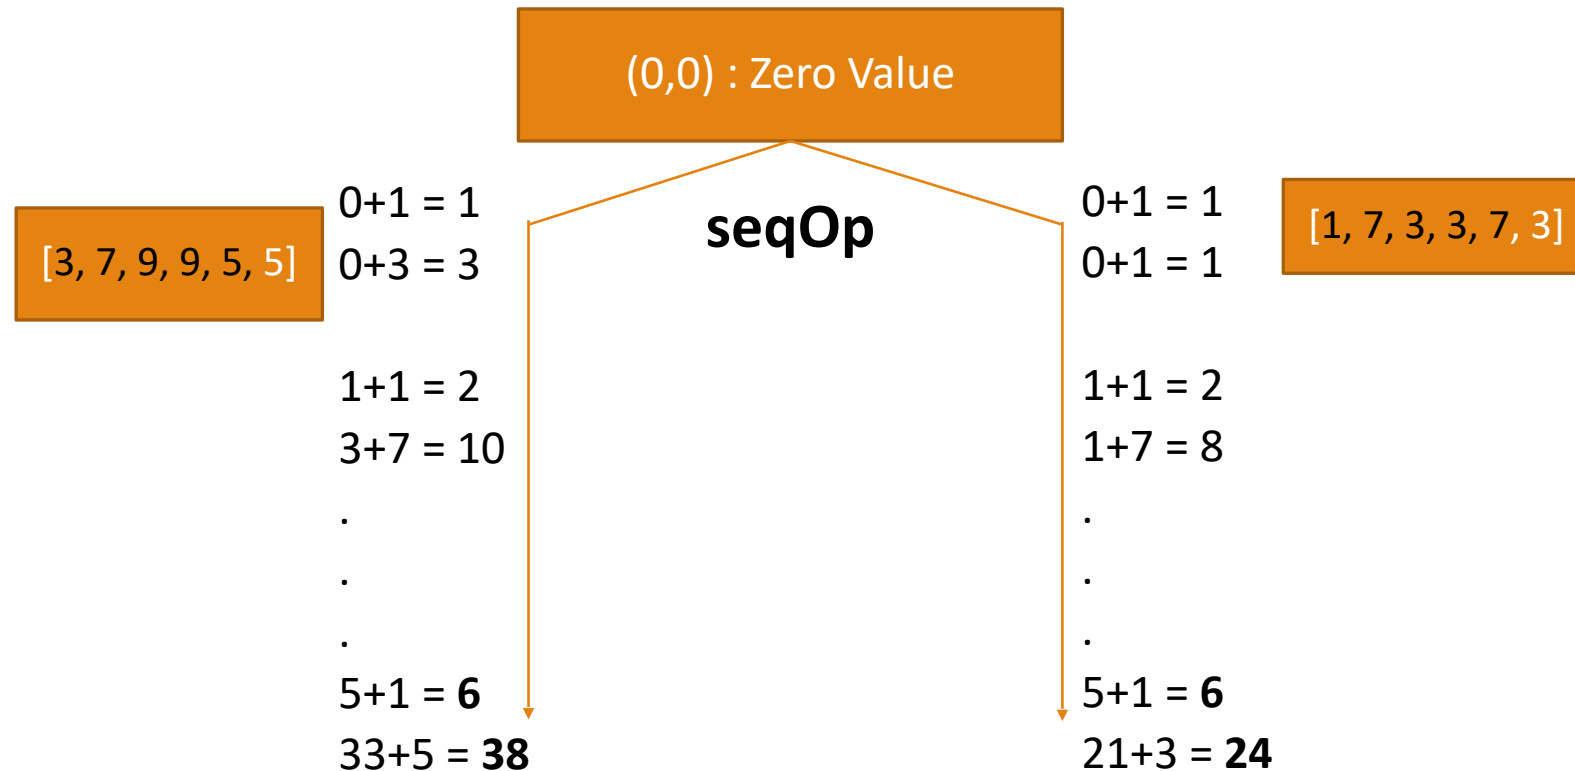
ex05

For data/numbers.txt, calculate the number and sum of the odd numbers as a pair using the aggregate function.



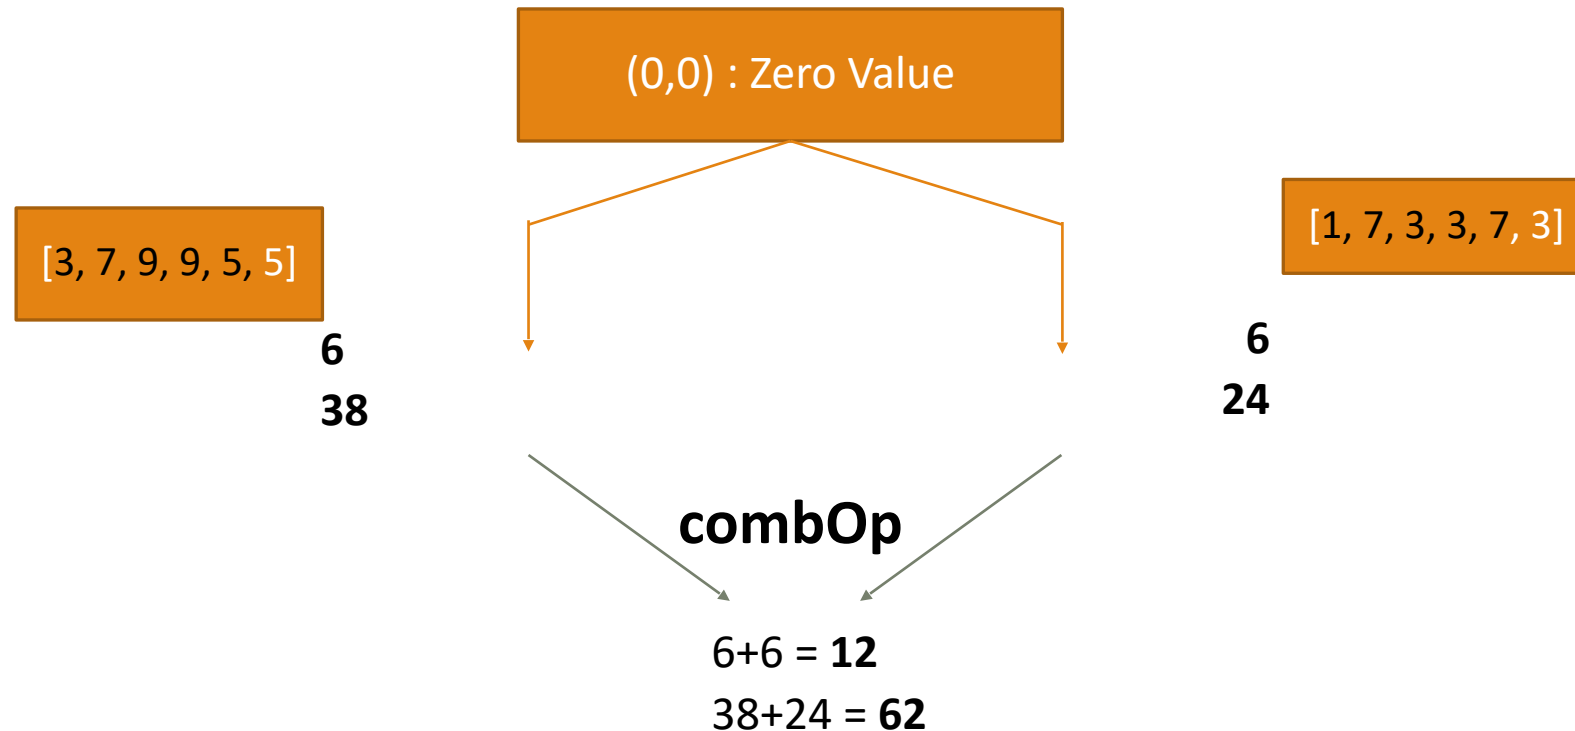
ex05

For data/numbers.txt, calculate the number and sum of the odd numbers as a pair using the aggregate function.



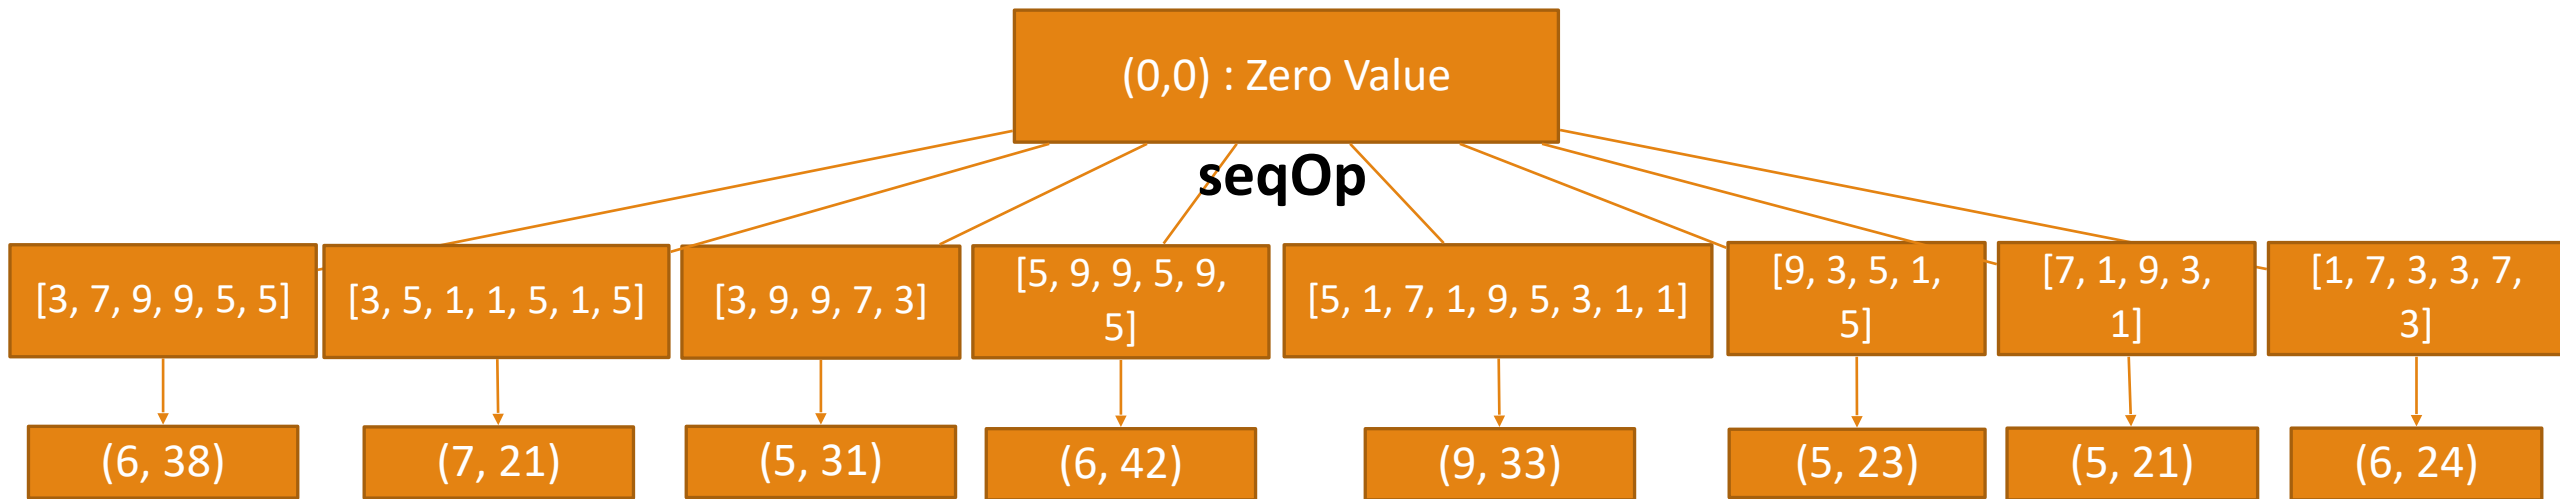
ex05

For data/numbers.txt, calculate the number and sum of the odd numbers as a pair using the aggregate function.



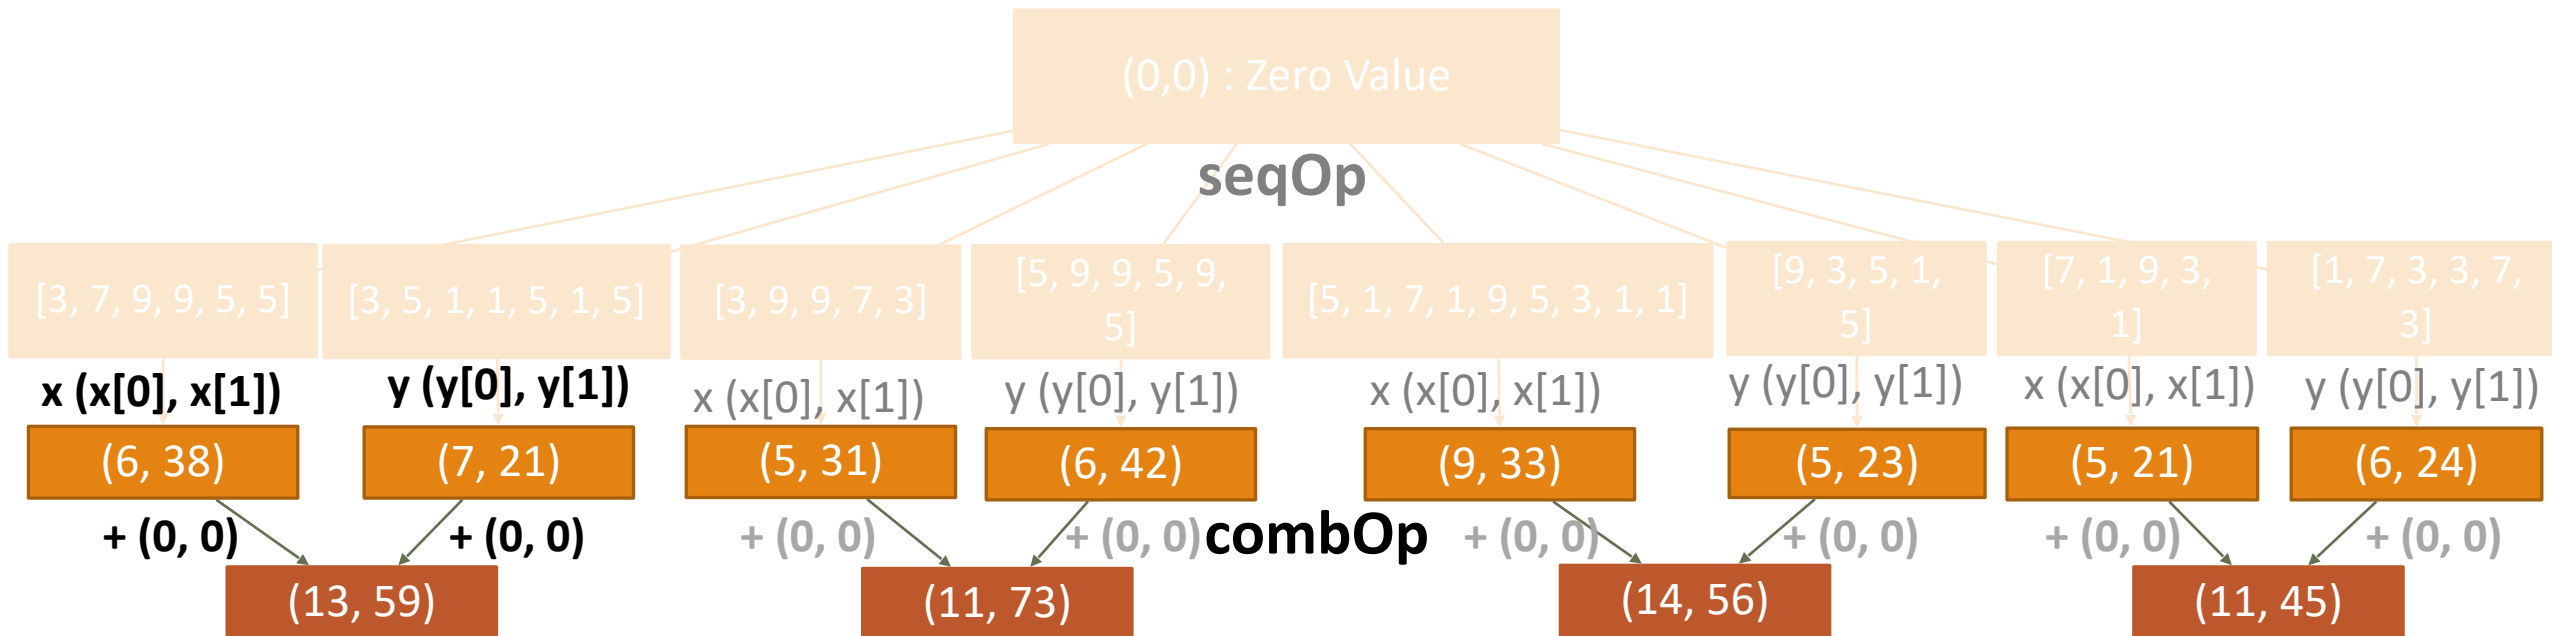
ex05

For data/numbers.txt, calculate the number and sum of the odd numbers as a pair using the aggregate function.



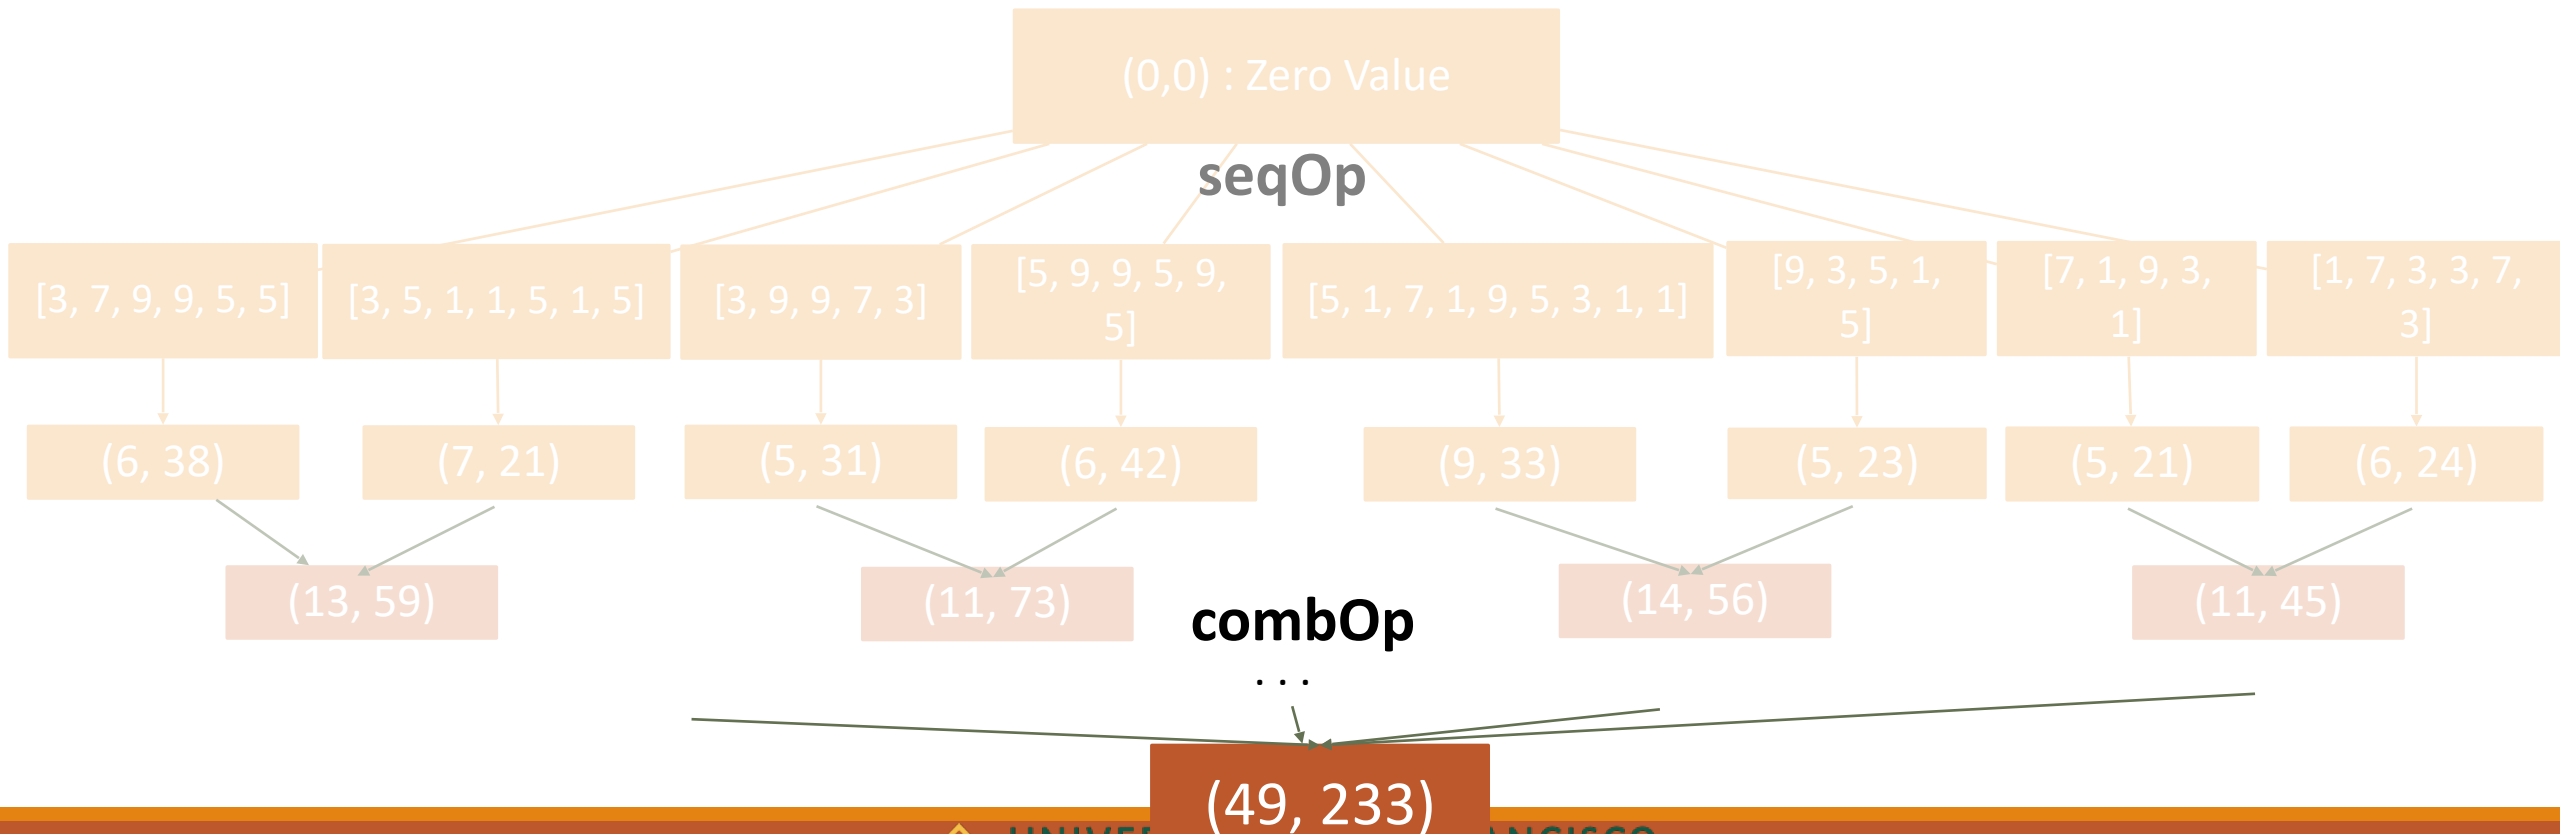
ex05

For data/numbers.txt, calculate the number and sum of the odd numbers as a pair using the aggregate function.



ex05

For data/numbers.txt, calculate the number and sum of the odd numbers as a pair using the aggregate function.



ex05

For data/numbers.txt, calculate the number and sum of the odd numbers as a pair using the aggregate function.

```
odd_numbers.aggregate((0, 0),  
                       (lambda x, y: (x[0] + 1, x[1] + y)),  
                       (lambda x, y: (x[0] + y[0], x[1] + y[1])))
```

(49, 233)



RDD Operations - Action

`saveAsTextFile(new_subdir_name)`

- Return multiple output files underneath the `new_subdir_name`, as Spark writes the output from multiple nodes.

<https://spark.apache.org/docs/latest/api/python/pyspark.html#pyspark.RDD>



UNIVERSITY OF SAN FRANCISCO

CHANGE THE WORLD FROM HERE

ex06

Store odd_numbers, the RDD in the “ex06_output” directory.

Read data from the “ex06_output” directory.

ex06

Store odd_numbers, the RDD in the “ex06_output” directory.

Read data from the “ex06_output” directory.

Store odd_numbers, the RDD in the “ex06_output” directory.

```
odd_numbers.saveAsTextFile("ex06_output")
```

Read data from the “ex06_output” directory.

```
odd_numbers_2 = sc.textFile("ex06_output")  
odd_numbers_2.collect()
```

Contents

RDD Operations

- Transformation
- Action





How many files are created in the output folder

1

A

8

B

9

C

None of the above

D



How many files are created in the output folder

1 **A**

8 **B**

9 **C**

None of the above **D**



How many files are created in the output folder

1 **A**

8 **B**

9 **C**

None of the above **D**

RDD Operations - Action

Action Operation Types

<code>collect()</code>	Return all the elements of the dataset as an array at the driver program. This is usually useful after a filter or other operation that returns a sufficiently small subset of the data.
<code>first()</code>	Return the first element of the dataset (similar to <code>take(1)</code>).
<code>take(n)</code>	Return an array with the first <code>n</code> elements of the dataset.
<code>count()</code>	Return the number of elements in the dataset.
<code>countByValue()</code>	Return the number of times each element occurs in the RDD.
<code>reduce(func)</code>	Combine the elements of the RDD together in parallel. (eg. Sum)
<code>fold(zero, func)</code>	Same as <code>reduce()</code> , but with the provided zero value.
<code>aggregate(zero, SeqOp, combOp)</code>	Similar to <code>reduce()</code> but used to return a different type.
<code>saveAsTextFile(path)</code>	Write the elements of the dataset as a text file (or set of text files) in a given directory in the local filesystem, HDFS or any other Hadoop-supported file system. Spark will call <code>toString</code> on each element to convert it to a line of text in the file.
<code>mean()</code> , <code>sum()</code> , <code>min()</code> , <code>max()</code> , <code>variance()</code> , <code>stdev()</code>	Returns mean, sum, minimum, maximum, variance and standard deviation.

Reference

Spark Online Documentation, <http://spark.apache.org/docs/latest/>

Zecevic, Petar, et al. Spark in Action, Manning, 2016.

Aven, Jeffrey. Apache Spark in 24 Hours, Sams Publishing , 2016.

Karau, Holden, et al. Learning spark: lightning-fast big data analysis. O'Reilly Media, Inc., 2015.

AWS Online Documentation, <https://aws.amazon.com/documentation/>

